



# Spring WebFlux

---

michael.schaffler@ciit.at



# 1. Introduction to Spring WebFlux

## What is Reactive Programming?

- non-blocking applications that are asynchronous and event-driven and require a small number of threads to scale vertically (i.e. within the JVM) rather than horizontally (i.e. through clustering)
- A key aspect of reactive applications is the concept of backpressure which is a mechanism to ensure producers don't overwhelm consumers
- Reactive programming also leads to a major shift from imperative to declarative async composition of logic. It is comparable to writing blocking code vs using the `CompletableFuture` from Java 8 to compose follow-up actions via lambda expressions

# Reactive API and Building Blocks



- Spring 5 uses Reactive Streams as the contract for communicating backpressure across async components and libraries
- Reactive Streams is a specification created through industry collaboration that has also been adopted in Java 9 as `java.util.concurrent.Flow`
- The Spring Framework uses Reactor internally for its own reactive support. Reactor is a Reactive Streams implementation that further extends the basic Reactive Streams Publisher contract with the Flux and Mono composable API types to provide declarative operations
- The Spring Framework exposes Flux and Mono in many of its own reactive APIs. At the application level Spring provides choice and fully supports the use of RxJava

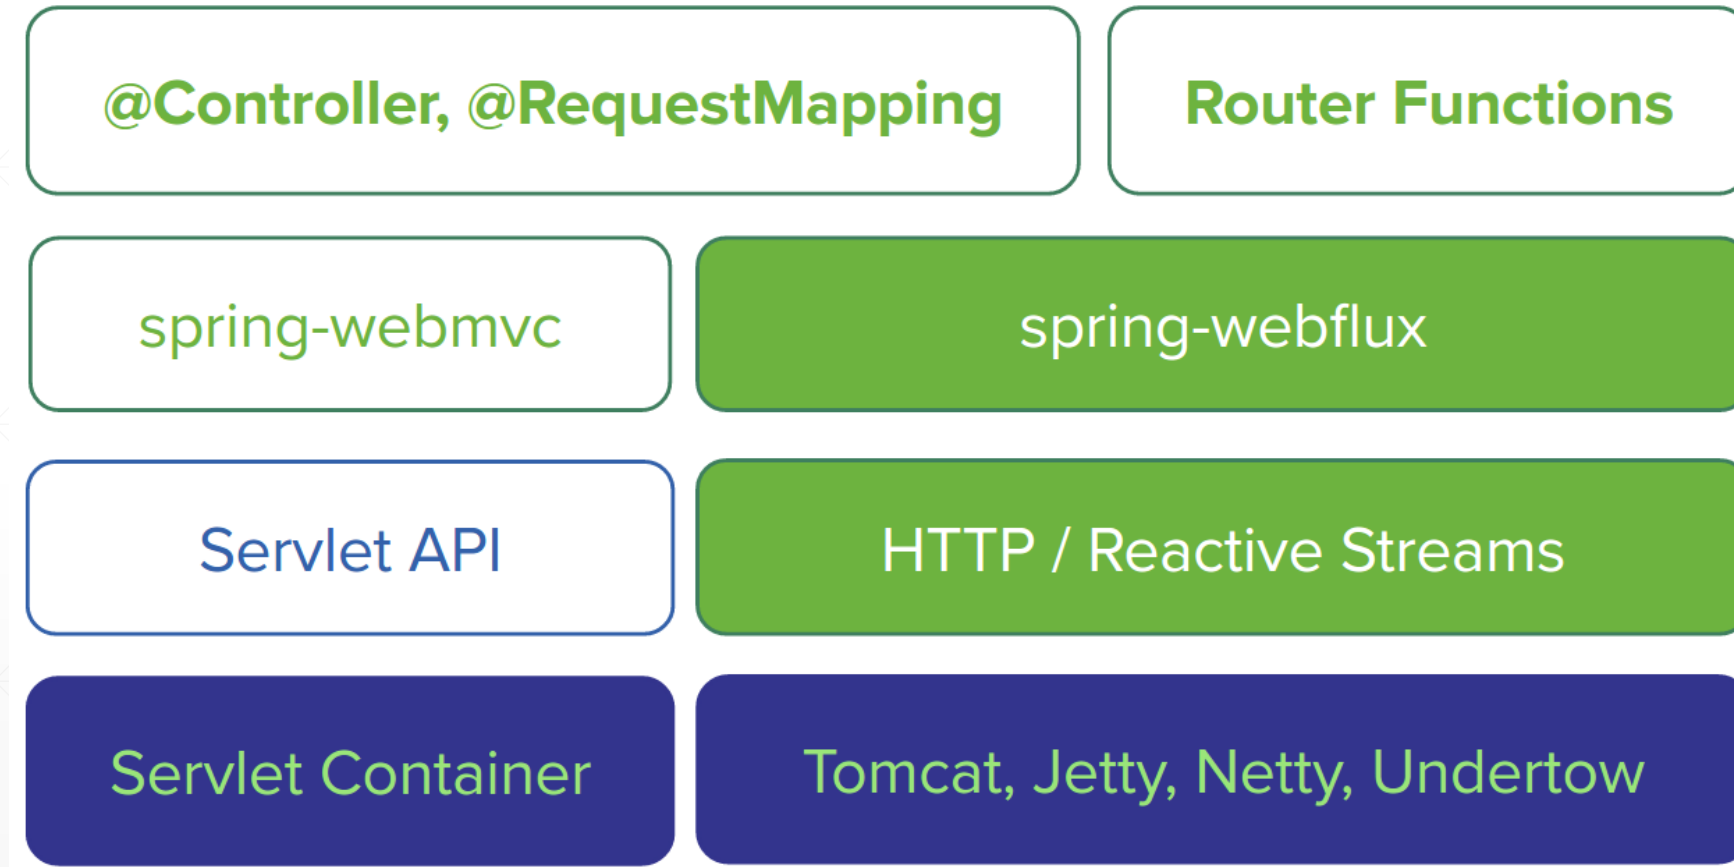
# Spring WebFlux Module



- Spring Framework 5 includes a new spring-webflux module. The module contains support for reactive HTTP and WebSocket clients as well as for reactive server web applications including REST, HTML browser, and WebSocket style interactions.
- **Server Side** - 2 distinct programming models
  - Annotation-based with @Controller and the other annotations supported also with Spring MVC
  - Functional, Java 8 lambda style routing and handling

Both programming models are executed on the same reactive foundation that adapts non-blocking HTTP runtimes to the Reactive Streams API

# Server side stack



# Annotation-based Programming Model



- The same `@Controller` programming model and the same annotations used in Spring MVC are also supported in WebFlux
- The main difference is that the underlying core, framework contracts — i.e. `HandlerMapping`, `HandlerAdapter`, are non-blocking and operate on the reactive `ServerHttpRequest` and `ServerHttpResponse`

# Annotation-based Programming Model (cont.)



```
@RestController
public class PersonController {

    private final PersonRepository repository;

    public PersonController(PersonRepository repository) {
        this.repository = repository;
    }

    @PostMapping("/person")
    Mono<Void> create(@RequestBody Publisher<Person> personStream) {
        return this.repository.save(personStream).then();
    }

    @GetMapping("/person")
    Flux<Person> list() {
        return this.repository.findAll();
    }

    @GetMapping("/person/{id}")
    Mono<Person> findById(@PathVariable String id) {
        return this.repository.findOne(id);
    }
}
```

# Functional Programming Model: HandlerFunctions



- Incoming HTTP requests are handled by a `HandlerFunction`, which is essentially a function that takes a `ServerRequest` and returns a `Mono<ServerResponse>`. The annotation counterpart to a handler function would be a method with `@RequestMapping`.
- `ServerRequest` gives access to various HTTP request elements: the method, URI, query parameters, and — through the separate `ServerRequest.Headers` interface — the headers. Access to the body is provided through the body methods
- Similarly, `ServerResponse` provides access to the HTTP response. Since it is immutable, you create a `ServerResponse` with a builder. The builder allows you to set the response status, add response headers, and provide a body





# HandlerFunctions

```
import static org.springframework.http.MediaType.APPLICATION_JSON;
import static org.springframework.web.reactive.function.BodyInserters.fromObject;

public class PersonHandler {

    private final PersonRepository repository;

    public PersonHandler(PersonRepository repository) {
        this.repository = repository;
    }

    public Mono<ServerResponse> listPeople(ServerRequest request) { ❶
        Flux<Person> people = repository.allPeople();
        return ServerResponse.ok().contentType(APPLICATION_JSON).body(people, Person.class);
    }

    public Mono<ServerResponse> createPerson(ServerRequest request) { ❷
        Mono<Person> person = request.bodyToMono(Person.class);
        return ServerResponse.ok().build(repository.savePerson(person));
    }

    public Mono<ServerResponse> getPerson(ServerRequest request) { ❸
        int personId = Integer.valueOf(request.pathVariable("id"));
        Mono<ServerResponse> notFound = ServerResponse.notFound().build();
        Mono<Person> personMono = this.repository.getPerson(personId);
        return personMono
            .then(person -> ServerResponse.ok().contentType(APPLICATION_JSON).body(fromObject(person)))
            .otherwiseIfEmpty(notFound);
    }
}
```

- ❶ `listPeople` is a handler function that returns all `Person` objects found in the repository as JSON.
- ❷ `createPerson` is a handler function that stores a new `Person` contained in the request body. Note that `PersonRepository.savePerson(Person)` returns `Mono<Void>`: an empty `Mono` that emits a completion signal when the person has been read from the request and stored. So we use the `build(Publisher<Void>)` method to send a response when that completion signal is received, i.e. when the `Person` has been saved.
- ❸ `getPerson` is a handler function that returns a single person, identified via the path variable `id`. We retrieve that `Person` via the repository, and create a JSON response if it is found. If it is not found, we use `otherwiseIfEmpty(Mono<T>)` to return a 404 Not Found response.

# Functional Programming Model: RouterFunctions



- Incoming requests are routed to handler functions with a RouterFunction, which is a function that takes a ServerRequest, and returns a Mono<HandlerFunction>
- If a request matches a particular route, a handler function is returned; otherwise it returns an empty Mono. The RouterFunction has a similar purpose as the @RequestMapping annotation in @Controller classes
- Typically, you do not write router functions yourself, but rather use RouterFunctions.route(RequestPredicate, HandlerFunction) to create one using a request predicate and handler function

# Functional Programming Model: RouterFunctions



```
import static org.springframework.http.MediaType.APPLICATION_JSON;
import static org.springframework.web.reactive.function.server.RequestPredicates.*;

PersonRepository repository = ...
PersonHandler handler = new PersonHandler(repository);

RouterFunction<ServerResponse> personRoute =
    route(GET("/person/{id}").and(accept(APPLICATION_JSON)), handler::getPerson)
        .andRoute(GET("/person").and(accept(APPLICATION_JSON)), handler::listPeople)
        .andRoute(POST("/person").and(contentType(APPLICATION_JSON)), handler::createPerson);
```

Besides router functions, you can also compose request predicates, by calling `RequestPredicate.and(RequestPredicate)` or `RequestPredicate.or(RequestPredicate)`. These work as expected: for `and` the resulting predicate matches if **both** given predicates match; `or` matches if **either** predicate does. Most of the predicates found in `RequestPredicates` are compositions. For instance, `RequestPredicates.GET(String)` is a composition of `RequestPredicates.method(HttpMethod)` and `RequestPredicates.path(String)`.

# Running a Server



```
RouterFunction<ServerResponse> route = ...
HttpHandler httpHandler = RouterFunctions.toHttpHandler(route);
HttpServlet servlet = new ServletHttpHandlerAdapter(httpHandler);
Tomcat server = new Tomcat();
Context rootContext = server.addContext("", System.getProperty("java.io.tmpdir"));
Tomcat.addServlet(rootContext, "servlet", servlet);
rootContext.addServletMapping("/", "servlet");
tomcatServer.start();
```

- The `HttpHandler` allows you to run on a wide variety of reactive runtimes: Reactor Netty, RxNetty, Servlet 3.1+, and Undertow

# HandlerFilterFunction



```
import static org.springframework.http.HttpStatus.UNAUTHORIZED;

SecurityManager securityManager = ...
RouterFunction<ServerResponse> route = ...

RouterFunction<ServerResponse> filteredRoute =
    route.filter(request, next) -> {
        if (securityManager.allowAccessTo(request.path())) {
            return next.handle(request);
        }
        else {
            return ServerResponse.status(UNAUTHORIZED).build();
        }
    });
```

# Client Side



- WebFlux includes a functional, reactive WebClient that offers a fully non-blocking and reactive alternative to the RestTemplate

```
WebClient client = WebClient.create("http://example.com");

Mono<Account> account = client.get()
    .url("/accounts/{id}", 1L)
    .accept(APPLICATION_JSON)
    .exchange(request)
    .then(response -> response.bodyToMono(Account.class));
```



The `AsyncRestTemplate` also supports non-blocking interactions. The main difference is it can't support non-blocking streaming, like for example [Twitter one](#), because fundamentally it's still based and relies on `InputStream` and `OutputStream`.

# Request and Response Body Conversion



- The spring-core module provides reactive Encoder and Decoder contracts that enable the serialization of a Flux of bytes to and from typed objects. The spring-web module adds JSON (Jackson) and XML (JAXB) implementations for use in web applications as well as others for SSE streaming and zero-copy file transfer.
- the request body can be one of the following way and it will be decoded automatically in both the annotation and the functional programming models:

- `Account account` — the account is deserialized without blocking before the controller is invoked.
- `Mono<Account> account` — the controller can use the `Mono` to declare logic to be executed after the account is deserialized.
- `Single<Account> account` — same as with `Mono` but using RxJava
- `Flux<Account> accounts` — input streaming scenario.
- `Observable<Account> accounts` — input streaming with RxJava.

# Request and Response Body Conversion (cont.)



The response body can be one of the following:

- `Mono<Account>` — serialize without blocking the given Account when the `Mono` completes.
- `Single<Account>` — same but using RxJava.
- `Flux<Account>` — streaming scenario, possibly SSE depending on the requested content type.
- `Observable<Account>` — same but using RxJava `Observable` type.
- `Flowable<Account>` — same but using RxJava 2 `Flowable` type.
- `Flux<ServerSentEvent>` — SSE streaming.
- `Mono<Void>` — request handling completes when the `Mono` completes.
- `Account` — serialize without blocking the given Account; implies a synchronous, non-blocking controller method.
- `void` — specific to the annotation-based programming model, request handling completes when the method returns;



# Reactive WebSocket Support



```
@Bean
public HandlerMapping websocketMapping() {
    Map<String, WebSocketHandler> map = new HashMap<>();
    map.put("/foo", new FooWebSocketHandler());
    map.put("/bar", new BarWebSocketHandler());

    SimpleUrlHandlerMapping mapping = new SimpleUrlHandlerMapping();
    mapping.setUrlMap(map);
    return mapping;
}

@Bean
public WebSocketHandlerAdapter handlerAdapter() {
    return new WebSocketHandlerAdapter();
}
```

On the client side create a `WebSocketClient` for one of the supported libraries listed above:

```
WebSocketClient client = new ReactorNettyWebSocketClient();
client.execute("ws://localhost:8080/echo", session -> {... }).blockMillis(5000);
```

# Testing



- The spring-test module includes a `WebTestClient` that can be used to test `WebFlux` server endpoints with or without a running server.
- Tests without a running server are comparable to `MockMvc` from Spring MVC where mock request and response are used instead of connecting over the network using a socket. The `WebTestClient` however can also perform tests against a running server

# Manual Bootstrapping



- For dependencies start with spring-webflux and spring-context. Then add jackson-databind and io.netty:netty-buffer for JSON support. Lastly add the dependencies for one of the supported runtimes, e.g. Tomcat, Jetty, Undertow...

For the **annotation-based programming model** bootstrap with:

```
ApplicationContext context = new AnnotationConfigApplicationContext(DelegatingWebFluxConfiguration.class); // (1)
HttpHandler handler = DispatcherHandler.toHttpHandler(context); // (2)
```

The above loads default Spring Web framework configuration (1), then creates a `DispatcherHandler`, the main class driving request processing (2), and adapts it to `HttpHandler` — the lowest level Spring abstraction for reactive HTTP request handling.

For the **functional programming model** bootstrap as follows:

```
ApplicationContext context = new AnnotationConfigApplicationContext(); // (1)
context.registerBean(FooBean.class, () -> new FooBeanImpl()); // (2)
context.registerBean(BarBean.class); // (3)

HttpHandler handler = WebHttpHandlerBuilder
    .webHandler(RouterFunctions.toHttpHandler(...))
    .applicationContext(context)
    .build(); // (4)
```

The above creates an `AnnotationConfigApplicationContext` instance (1) that can take advantage of the new functional bean registration API (2) to register beans using a Java 8 `Supplier` or just by specifying its class (3). The `HttpHandler` is created using `WebHttpHandlerBuilder` (4).



# Manual Bootstrapping



- For dependencies start with spring-webflux and spring-context. Then add jackson-databind and io.netty:netty-buffer for JSON support. Lastly add the dependencies for one of the supported runtimes, e.g. Tomcat, Jetty, Undertow...

For the **annotation-based programming model** bootstrap with:

```
ApplicationContext context = new AnnotationConfigApplicationContext(DelegatingWebFluxConfiguration.class); // (1)
HttpHandler handler = DispatcherHandler.toHttpHandler(context); // (2)
```

The above loads default Spring Web framework configuration (1), then creates a `DispatcherHandler`, the main class driving request processing (2), and adapts it to `HttpHandler` — the lowest level Spring abstraction for reactive HTTP request handling.

For the **functional programming model** bootstrap as follows:

```
ApplicationContext context = new AnnotationConfigApplicationContext(); // (1)
context.registerBean(FooBean.class, () -> new FooBeanImpl()); // (2)
context.registerBean(BarBean.class); // (3)

HttpHandler handler = WebHttpHandlerBuilder
    .webHandler(RouterFunctions.toHttpHandler(...))
    .applicationContext(context)
    .build(); // (4)
```

The above creates an `AnnotationConfigApplicationContext` instance (1) that can take advantage of the new functional bean registration API (2) to register beans using a Java 8 `Supplier` or just by specifying its class (3). The `HttpHandler` is created using `WebHttpHandlerBuilder` (4).



# Manual Bootstrapping



- For dependencies start with spring-webflux and spring-context. Then add jackson-databind and io.netty:netty-buffer for JSON support. Lastly add the dependencies for one of the supported runtimes, e.g. Tomcat, Jetty, Undertow...

For the **annotation-based programming model** bootstrap with:

```
ApplicationContext context = new AnnotationConfigApplicationContext(DelegatingWebFluxConfiguration.class); // (1)
HttpHandler handler = DispatcherHandler.toHttpHandler(context); // (2)
```

The above loads default Spring Web framework configuration (1), then creates a `DispatcherHandler`, the main class driving request processing (2), and adapts it to `HttpHandler` — the lowest level Spring abstraction for reactive HTTP request handling.

For the **functional programming model** bootstrap as follows:

```
ApplicationContext context = new AnnotationConfigApplicationContext(); // (1)
context.registerBean(FooBean.class, () -> new FooBeanImpl()); // (2)
context.registerBean(BarBean.class); // (3)

HttpHandler handler = WebHttpHandlerBuilder
    .webHandler(RouterFunctions.toHttpHandler(...))
    .applicationContext(context)
    .build(); // (4)
```

The above creates an `AnnotationConfigApplicationContext` instance (1) that can take advantage of the new functional bean registration API (2) to register beans using a Java 8 `Supplier` or just by specifying its class (3). The `HttpHandler` is created using `WebHttpHandlerBuilder` (4).





# Spring WebFlux

---

michael.schaffler@ciit.at